

LAB 3

Introduction

In this laboratory we are going to study the implementation of "divide and conquer" strategy using OpneMP paralization. **(ha de continuar)**

Task decomposition analysis for Mergesort

Divide and conquer

Task decomposition analysis with Tareador

```

void merge(long n, T left[n], T right[n], T result[n*2], long start, long
length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start,
length);
    } else {
        // Recursive decomposition
        tareador_start_task("merge 0");
        merge(n, left, right, result, start,
length/2);

        tareador_end_task("merge 0");
        tareador_start_task("merge 1");
        merge(n, left, right, result, start +
length/2, length/2);

        tareador_end_task("merge 1");
    }
}

void multisort(long n, T data[n], T tmp[n]) {

    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        tareador_start_task("multisort 0");
        multisort(n/4L, &data[0], &tmp[0]);
        tareador_end_task("multisort 0");
        tareador_start_task("multisort 1");
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        tareador_end_task("multisort 1");
        tareador_start_task("multisort 2");
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        tareador_end_task("multisort 2");
        tareador_start_task("multisort 3");
        multisort(n/4L, &data[3L*n/4L],
&tmp[3L*n/4L]);

        tareador_end_task("multisort 3");

        tareador_start_task("merge_multi 0");
        merge(n/4L, &data[0], &data[n/4L], &tmp[0],
0, n/2L);

        tareador_end_task("merge_multi 0");
        tareador_start_task("merge_multi 1");
        merge(n/4L, &data[n/2L], &data[3L*n/4L],
&tmp[n/2L], 0, n/2L);

        tareador_end_task("merge_multi 1");
    }
}

```

```

    0, n);

    tareador_start_task("merge_multi 2");
    merge(n/2L, &tmp[0], &tmp[n/2L], &data[0],

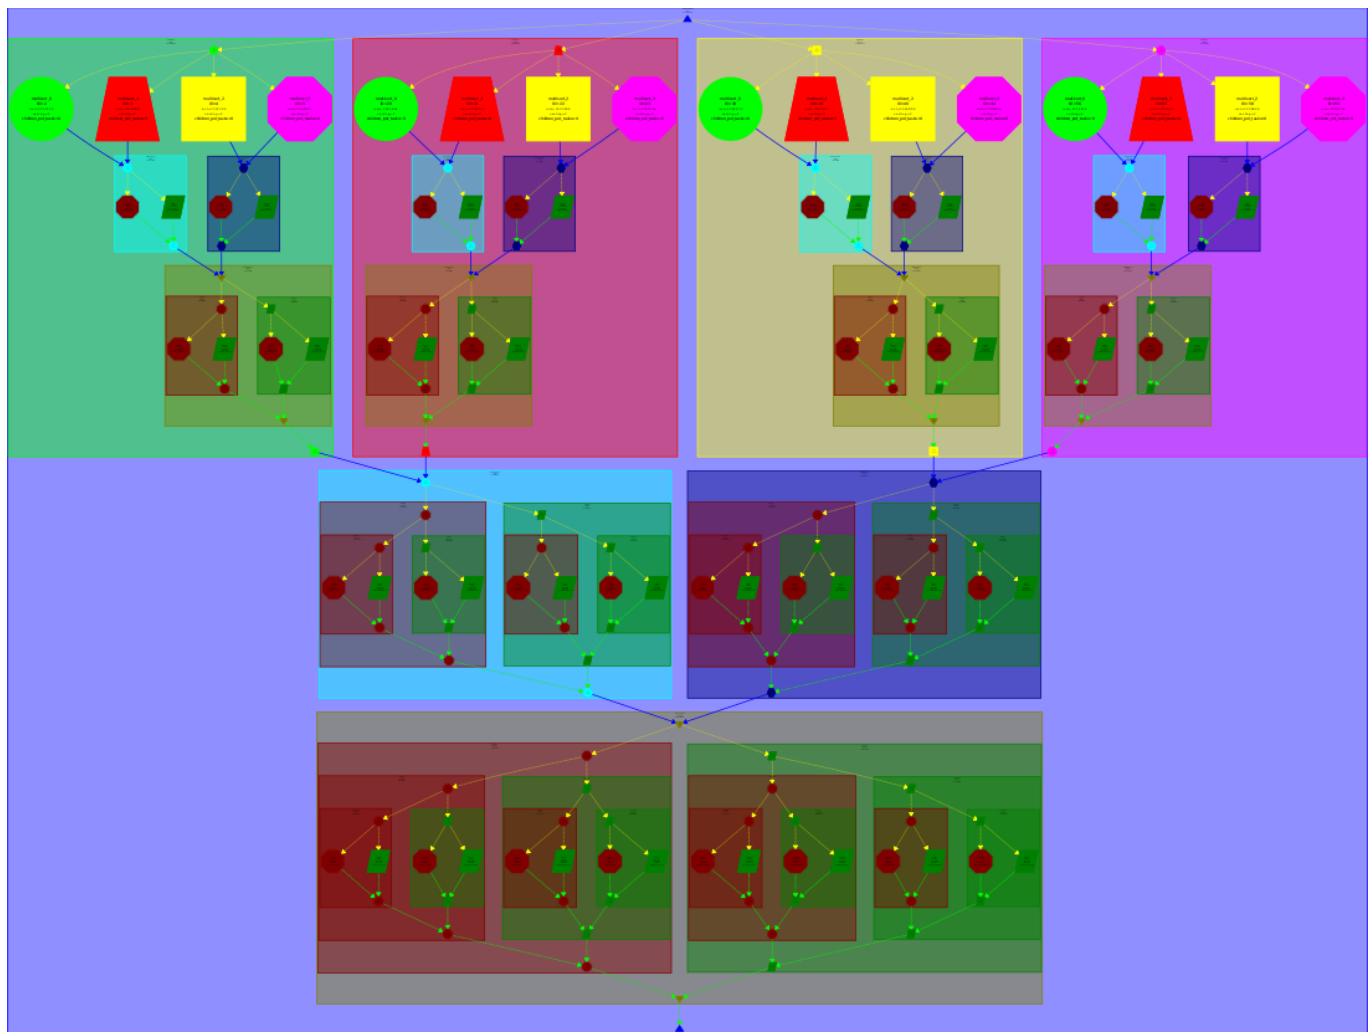
    tareador_end_task("merge_multi 2");

    } else {
        // Base case
        basicsort(n, data);
    }
}


```

[multisort-tareador.c](#)


That piece of code give us the following dependence graph:




We have added a tareador task in each of the recursive tasks, to fully visualize the possible parallelizations of the code.




Trace of multisort-tareador using 1 core



Trace of multisort-tareador using 2 core




Trace of multisort-tareador using 4 core




Trace of multisort-tareador using 8 core



Trace of multisort-tareador using 16 core







Trace of multisort-tareador using 64 core

| NUM CPUs | time (ns) | SPEED UP | EFFICIENCY |
|----------|-------------|------------------|-------------------|
| 1 | 20334411001 | 1 | 0.999360066616823 |
| 2 | 10173716001 | 1.99872013323365 | 0.99938619627572 |
| 4 | 5086725001 | 3.99754478510288 | 0.99938619627572 |
| 8 | 2550595001 | 7.97241858979085 | 0.996552323723856 |
| 16 | 1289922001 | 15.7640624667507 | 0.985253904171916 |
| 32 | 1289909001 | 15.7642213406029 | 0.492631916893841 |
| 64 | 1289909001 | 15.7642213406029 | 0.246315958446921 |

The following pots represents number of cores vs. plot of time, speed up and efficiency, respectively.



Number of cores vs. execution time plot and number of cores vs. speed up plot.



Number of cores vs. efficiency plot.

The multisort program parallizes ideally untill 16 cores, when , the dependences between the multisort marge calls, as seen in the tareador capture of the tasks, doesn't allow for more. The efficiency drops going further than 16 threads (i.e fig: efficiency plot and 32,64 captures) Proving that adding more CPUs is pointless.

Shared-memory parallelization with OpenMP tasks

Task cut-off mechanism

First of all we are going to implement two versions of multisort algorithm. One using a tree strategy and another using leaf strategy.

Tree strategy code:

```
void merge(long n, T left[n], T right[n], T result[n*2], long start, long
length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        #pragma omp taskgroup
        {
            #pragma omp task
            merge(n, left, right, result, start, length/2);
            #pragma omp task
            merge(n, left, right, result, start + length/2, length/2);
        }
    }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L)
    {
        #pragma omp taskgroup
        {
            #pragma omp task
            multisort(n/4L, &data[0], &tmp[0]);
            #pragma omp task
            multisort(n/4L, &data[n/4L], &tmp[n/4L]);
            #pragma omp task
            multisort(n/4L, &data[n/2L], &tmp[n/2L]);
            #pragma omp task
            multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);
        }
        #pragma omp taskgroup
        {
            #pragma omp task
            merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
            #pragma omp task
            merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);
            // #pragma omp taskwait
        }
        #pragma omp task
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
    }
}
```



```

    else
    {
        basicsort(n, data);
    }
}

```

multisort-omp-tree.c

| | multisort | merge |
|----------------|-----------|----------|
| THREAD 1.1.1 | 266 | 3,098 |
| THREAD 1.1.2 | 259 | 2,956 |
| THREAD 1.1.3 | 329 | 1,990 |
| THREAD 1.1.4 | 280 | 2,782 |
| THREAD 1.1.5 | 329 | 2,028 |
| THREAD 1.1.6 | 294 | 1,812 |
| THREAD 1.1.7 | 301 | 1,872 |
| THREAD 1.1.8 | 329 | 1,896 |
| | | |
| Total | 2,387 | 18,434 |
| Average | 298.38 | 2,304.25 |
| Maximum | 329 | 3,098 |
| Minimum | 259 | 1,812 |
| StDev | 26.87 | 506.70 |
| Avg/Max | 0.91 | 0.74 |

Paraver capture of number of tasks created using tree strategy

Lear strategy code:

```

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        multisort(n/4L, &data[0], &tmp[0]);
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);

        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0,
n/2L);

        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
    }
    else {
        // Base case
        #pragma omp taskgroup
        {
            #pragma omp task
            basicsort(n, data);
        }
        //#pragma omp taskwait
    }
}

```

multisort-omp-leaf.c

| | multisort |
|----------------|-----------|
| THREAD 1.1.1 | 136 |
| THREAD 1.1.2 | 102 |
| THREAD 1.1.3 | 100 |
| THREAD 1.1.4 | 167 |
| THREAD 1.1.5 | 147 |
| THREAD 1.1.6 | 104 |
| THREAD 1.1.7 | 158 |
| THREAD 1.1.8 | 110 |
| | |
| Total | 1,024 |
| Average | 128 |
| Maximum | 167 |
| Minimum | 100 |
| StDev | 25.51 |
| Avg/Max | 0.77 |

Paraver capture of number of tasks created using leaf strategy

For a cut-off strategy we need to marge tree and leaf methods. Until a specified number of created tasks the program will create tasks like tree strategy, recursively. Then, when the limit is reached, the program is going to create sequentially one task for each leaf of leaf task.

Cut-off strategy code:

```
void merge(long n, T left[n], T right[n], T result[n*2], long start, long
length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        if(NUM_TASKS < CUTOFF){
            #pragma omp taskgroup
            {
                NUM_TASKS++;
                #pragma omp task
                merge(n, left, right,
result, start, length/2);

                NUM_TASKS++;
                #pragma omp task
                merge(n, left, right,
result, start + length/2, length/2);
            }
            NUM_TASKS -= 2;
        }
        else{
            merge(n, left, right, result,
```

```

start, length/2);
                                merge(n, left, right, result, start
+ length/2, length/2);
                                }
                                }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        if(NUM_TASKS < CUTOFF){
            #pragma omp taskgroup
            {
                NUM_TASKS++;
                #pragma omp task
                multisort(n/4L, &data[0],
&tmp[0]);

                NUM_TASKS++;
                #pragma omp task
                multisort(n/4L,
&data[n/4L], &tmp[n/4L]);

                NUM_TASKS++;
                #pragma omp task
                multisort(n/4L,
&data[n/2L], &tmp[n/2L]);

                NUM_TASKS++;
                #pragma omp task
                multisort(n/4L,
&data[3L*n/4L], &tmp[3L*n/4L]);
            }
            NUM_TASKS -= 4;
            #pragma omp taskgroup
            {
                NUM_TASKS++;
                #pragma omp task
                merge(n/4L, &data[0],
&data[n/4L], &tmp[0], 0, n/2L);

                NUM_TASKS++;
                #pragma omp task
                merge(n/4L, &data[n/2L],
&data[3L*n/4L], &tmp[n/2L], 0, n/2L);

                // #pragma omp taskwait
            }
            NUM_TASKS -= 2;

            NUM_TASKS++;
            #pragma omp task
            merge(n/2L, &tmp[0], &tmp[n/2L],
&data[0], 0, n);
        }
    }
    else{
        // Recursive decomposition
        multisort(n/4L, &data[0], &tmp[0]);
        multisort(n/4L, &data[n/4L],

```

```

&tmp[n/4L]);
                                multisort(n/4L, &data[n/2L],
&tmp[n/2L]);
                                multisort(n/4L, &data[3L*n/4L],
&tmp[3L*n/4L]);
                                merge(n/4L, &data[0], &data[n/4L],
&tmp[0], 0, n/2L);
                                merge(n/4L, &data[n/2L],
&data[3L*n/4L], &tmp[n/2L], 0, n/2L);
                                merge(n/2L, &tmp[0], &tmp[n/2L],
&data[0], 0, n);
                                }
    } else {
        if(NUM_TASKS >= CUTOFF){
            #pragma omp task
            basicsort(n, data);
            #pragma omp taskwait
        }
        else{
            basicsort(n, data);
        }
    }
}


```

multisort-omp-cutoff.c


| | multisort | merge |
|--------------|-----------|-------|
| THREAD 1.1.1 | 263 | - |
| THREAD 1.1.2 | 256 | - |
| THREAD 1.1.3 | 256 | 2 |
| THREAD 1.1.4 | 256 | 2 |
| THREAD 1.1.5 | - | - |
| THREAD 1.1.6 | - | - |
| THREAD 1.1.7 | - | - |
| THREAD 1.1.8 | - | 2 |
| Total | 1,031 | 6 |
| Average | 257.75 | 2 |
| Maximum | 263 | 2 |
| Minimum | 256 | 2 |
| StDev | 3.03 | 0 |
| Avg/Max | 0.98 | 1 |

Paraver capture of number of tasks created using cut-off strategy

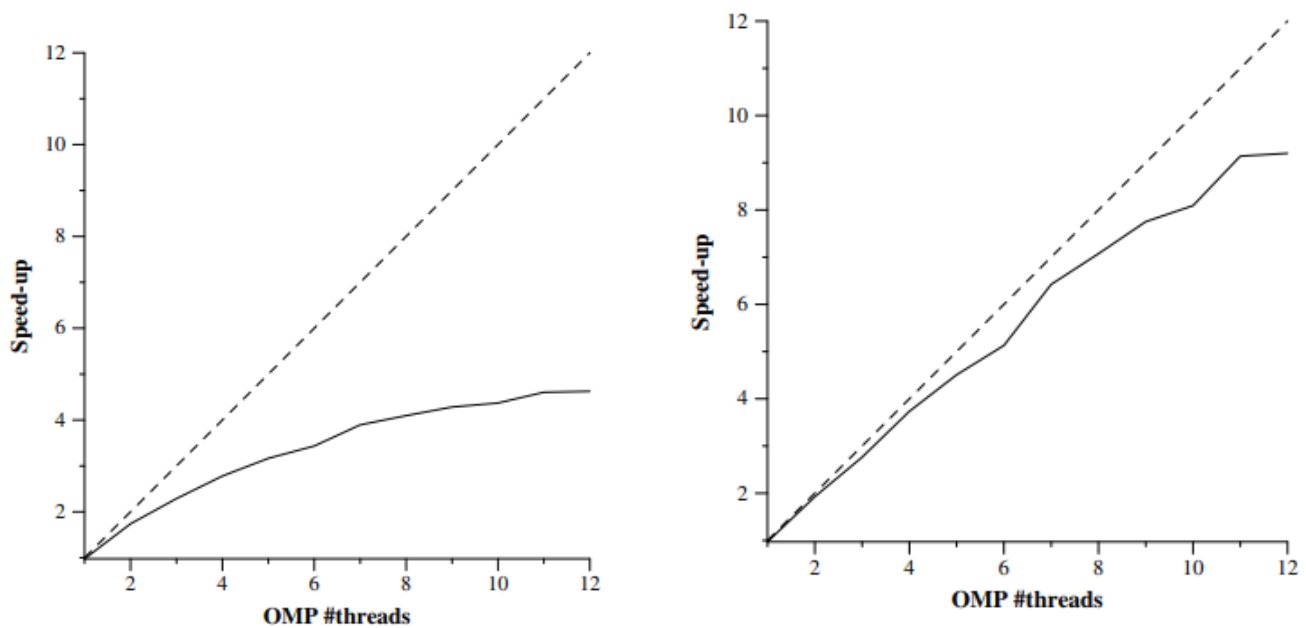
Donw bellow we are going to study the strong scalability of the three diferent versions of the multisort algorithm using cores vs time/speed up plots.



Elapsed time plot and speed up plot for tree version of multisort algorithm



Elapsed time plot and speed up plot for leaf version of multisort algorithm



PostScript plots generated by submit-cutoff.sh

We can observe that the speed up at first plot (using the execution time of the whole program) reach a limit 10 cores equal than the second plot (only using the time of multisort function). Nevertheless the scalability ratio is better on the second plot, because is where the parallelization makes higher work.

Using OpenMP task dependencies

At last, we are going to implement, based in tree strategy, a task dependence model.

```
void merge(long n, T left[n], T right[n], T result[n*2], long start, long
length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        #pragma omp taskgroup
        {
            #pragma omp task
            merge(n, left, right, result,
start, length/2);

            #pragma omp task
            merge(n, left, right, result, start
+ length/2, length/2);
        }
    }
}

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        #pragma omp taskgroup
        {
            #pragma omp task
            multisort(n/4L, &data[0], &tmp[0]);
            #pragma omp task
            multisort(n/4L, &data[n/4L],
&tmp[n/4L]);

            #pragma omp task
            multisort(n/4L, &data[n/2L],
&tmp[n/2L]);

            #pragma omp task
            multisort(n/4L, &data[3L*n/4L],
&tmp[3L*n/4L]);
        }
        #pragma omp taskgroup
        {
            #pragma omp task
            merge(n/4L, &data[0], &data[n/4L],
&tmp[0], 0, n/2L);

            #pragma omp task
            merge(n/4L, &data[n/2L],
&data[3L*n/4L], &tmp[n/2L], 0, n/2L);
        }
        // #pragma omp taskwait

        #pragma omp task
    }
}
```

```
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
    } else {
        // Base case
        basicsort(n, data);
    }
}
```

[multisort-omp-dependences.c](#)

Conclusions

- Acabar explicar coses s2
 - explicar cutoff
- Fer s3
 - tot
 - opcionals
- Fer conclusió
- Acabar intro